

Software Technology for Internet of Things

Stream Processing

Aslak Johansen asjo@mmmi.sdu.dk

Apr 1, 2025

Part 0: Introduction

Anatomy of a (Typical) IoT Cloud


Key properties include:

- ▶ Large total number of incoming data streams.
- ▶ A large number of customers
 - ▶ Customers expect their data to stay private (legal reasons).
 - ▶ The processing setup is the same for all customers*.
- ▶ Processing have different justifications:
 - ▶ Required for promoted functionality.
 - ▶ Real-world evaluation of product line.
 - ▶ Pattern detection (e.g., for future products).
 - ▶ Health monitoring.
- ▶ A small number of the processing products are presented to the customer.
- ▶ Potentially large number of outgoing data streams.

* at least it is close.

Processing Graph

These have various names, including: graph processing, stream processing.



General Processing Case

A stream process is independent piece of logic associated with:

- ▶ State.
- ▶ Zero or more input streams.
- ▶ Zero or more output streams.

A **node** consumes input(s) in order to produce output(s).

Specific cases:

- ▶ **Source:** A node only producing output stream(s).
- ▶ **Sink:** A node only consuming input stream(s).

Specific Processing Properties

- ▶ **Flow Control** What is the right policy when data is consumed at a lower rate than it is produced?
 - ▶ Buffering or dropping?
- ▶ **Consumability** Does processing consume the data or should it be retained?
- ▶ **Trigger Condition** If a node consumes more than one input, when does it fire?
 - ▶ Synchronize inputs using a barrier or use latest values?
 - ▶ What happens if a value is dropped?
- ▶ **Trigger Arbitrariness** Can a node fire without inputs?
 - ▶ E.g., to implement timeouts.
- ▶ **Batch Size** Can we accept the latency penalty of operating in batches, and how big should we make them?
- ▶ **Fault Tolerance** What is the right policy when a node dies?

Concerns and Remedies

There are a few general problems when dealing with lots of concurrent pieces of logic:



- ▶ Context switching
 - ▶ Switching – on OS level – from the context of one process to the context of another is one of the most expensive operations.
 - ▶ **Preempting** a process thread with another is still an expensive operation.
 - ▶ **Multiplexing** can be *relatively* inexpensive.
- ▶ Side-Effects
 - ▶ Code with side-effects generate cache misses.
 - ▶ Without side-effects, concurrency is trivial[†].
- ▶ Network
 - ▶ In a **share-nothing architecture**, an update is satisfied by a(ny) single node.

"Going to disk is 25 million times slower than hitting a general purpose register. Design accordingly."



– Robert Love, GUADEC 2005

Part 1: System Design

Overview



Overview



Part 2: Message Queues

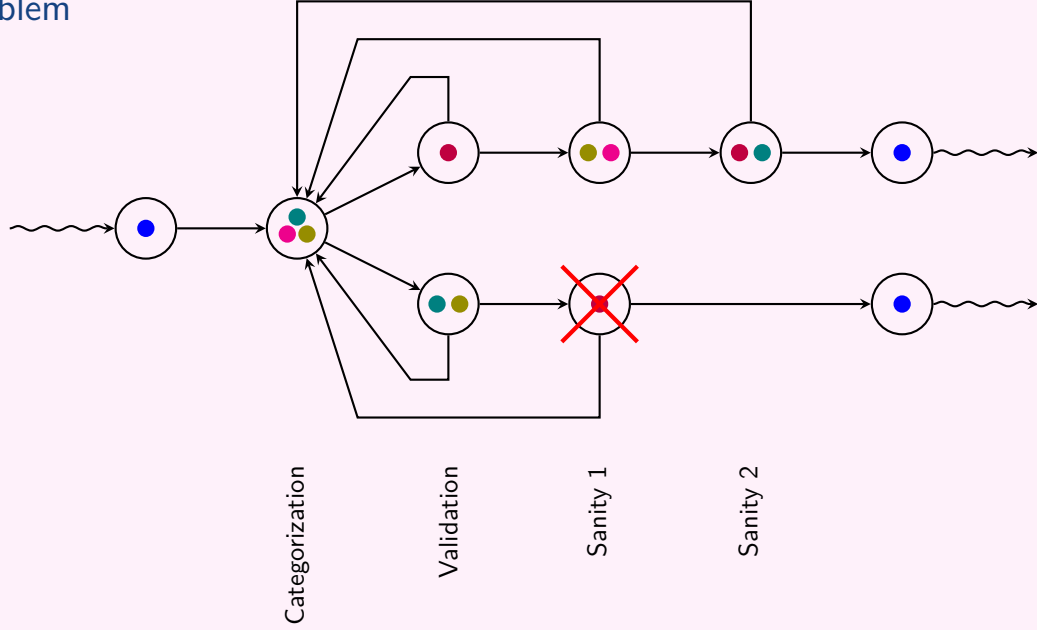
Case

Lets look at a **data ingestion** situation:

- ▶ Incoming data needs to be stored in a database.
- ▶ A few steps are inserted along this path covering
 - ▶ Categorization
 - ▶ Validation
 - ▶ Sanity checking
- ▶ A failed step may require another processing attempt (e.g., using the notion of a TTL).
- ▶ All incoming data needs to be processed, independently of the availability of the components along the way.

What could be defining for the design?

Problem



Part 3:

Event Busses


Case

Lets look at a **clock signal**:

- ▶ A stream of timestamps are being received.
- ▶ These needs to be adjusted for timezone.
- ▶ Timeliness matters.

What could be defining for the design?

Problem



Part 4:

Publish Subscribe Pattern

Publish Subscribe

Can someone describe what publish subscribe is?

A communication pattern, in which:

- ▶ Messages can be published through some **name** on a *broker*.
- ▶ Clients can subscribe to **names** on a broker.
- ▶ Whenever a message is published, all clients subscribing to that **name** are notified (i.e., forwarded the message).

Often referred to as *pubsub*.

Publish Subscribe

Publish Subscribe

Producer 1



Publish Subscribe

Producer 1

Producer 2

Publish Subscribe

Producer 1

Producer 2

Producer 3

Publish Subscribe

Producer 1

Producer 2

Producer 3

Producer 4

Publish Subscribe

Producer 1

Producer 2

Producer 3

Producer 4

Consumer 1

Publish Subscribe

Producer 1

Producer 2

Producer 3

Producer 4

Consumer 1

Consumer 2

Publish Subscribe

Producer 1

Producer 2

Producer 3


Producer 4

Consumer 1


Consumer 2

Consumer 3


Publish Subscribe




Publish Subscribe




Publish Subscribe



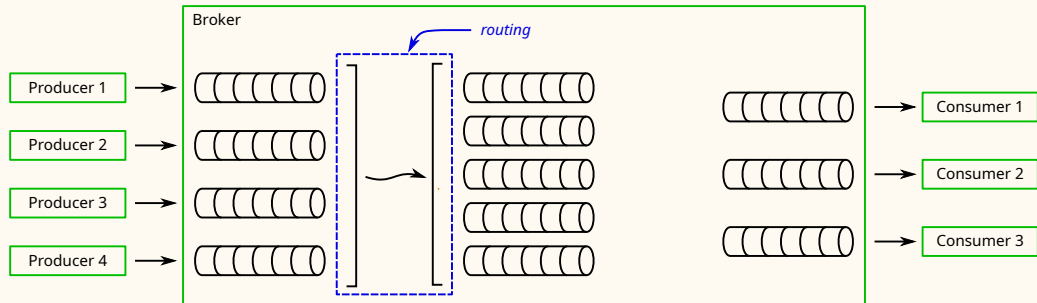
Publish Subscribe




Publish Subscribe




Publish Subscribe



Publish Subscribe



Publish Subscribe



Publish Subscribe Smart and Dumb Components

Dumb Broker
&
Smart Client

Smart Broker
&
Dumb Client




Publish Subscribe Smart and Dumb Components

Dumb Broker
&
Smart Client

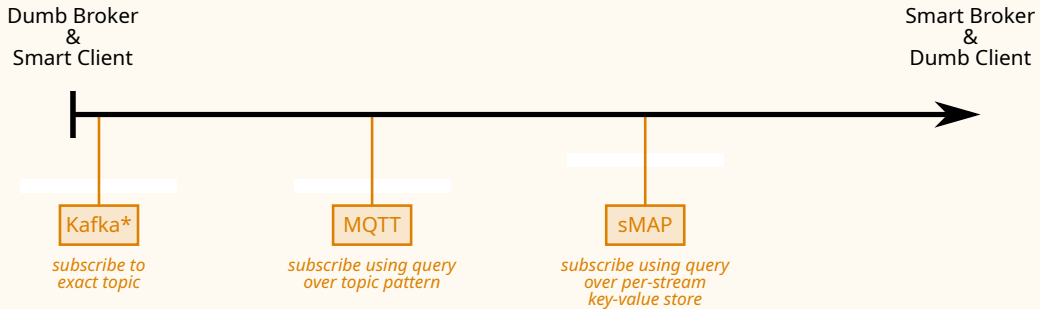
Smart Broker
&
Dumb Client



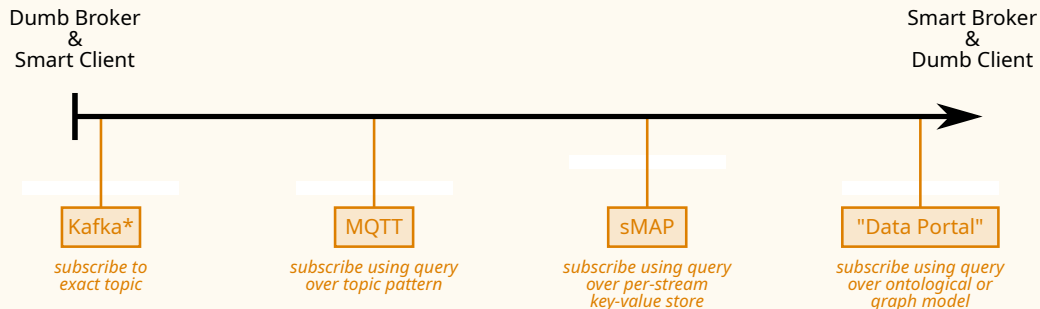
Publish Subscribe Smart and Dumb Components



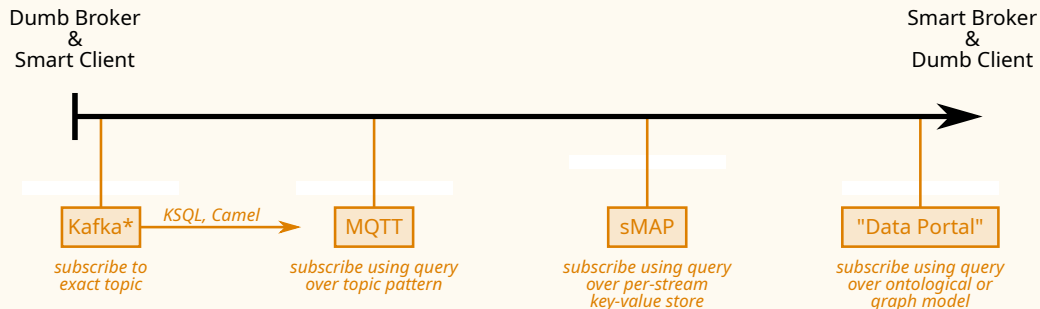
Publish Subscribe Smart and Dumb Components



Publish Subscribe Smart and Dumb Components



Publish Subscribe Smart and Dumb Components



Message Queuing Telemetry Transport (MQTT)

An example of a subsub protocol that is frequently used in the IoT space.

Properties:

- ▶ Subscriptions through topic patterns.
- ▶ Identifier patterns may be expressed using two kinds of wildcards:
 - ▶ Topics are matched against patterns.
 - ▶ Topics are structured strings and use "/" as separator.
 - ▶ The "+" wildcard matches a single level.
 - ▶ The "#" wildcard may match multiple levels.
- ▶ Quality of service (for each of the two transmissions): at-most-once, at-least-once and exactly-once.
- ▶ Last will and testament: On-disconnect events.
- ▶ Persistent sessions: Per-client offline buffering.
- ▶ Very simple protocol!
 - ▶ Broker and framework choice.
 - ▶ Innovation.

Part 5: Distributed Concerns

Dealing with System Failures

Rule: *“As the number of machines in a system goes towards infinity the time to one of them breaking goes to zero”*.

Accordingly, as we grow our system, the need for dealing with failures increases.

Several support systems have been designed to deal with this **on a service level**.

They generally:

- ▶ Can be told to run a service on n machines.
- ▶ Makes sure to restart a service when it goes down.
- ▶ Hosts distributed filesystem(s) on the same hardware as the services.
- ▶ Allows for geographical redundancy.

For **heterogeneous and low-granularity** deployments, Kubernetes and Ceph comes into play, with Docker Swarm being another option. For **homogeneous and high-granularity** deployment BEAM is king. Often these are combined.

Questions?

